

The background features abstract, flowing waves in shades of red, orange, and yellow, creating a dynamic and modern aesthetic.

INTRODUCTION TO ORCHESTRATION

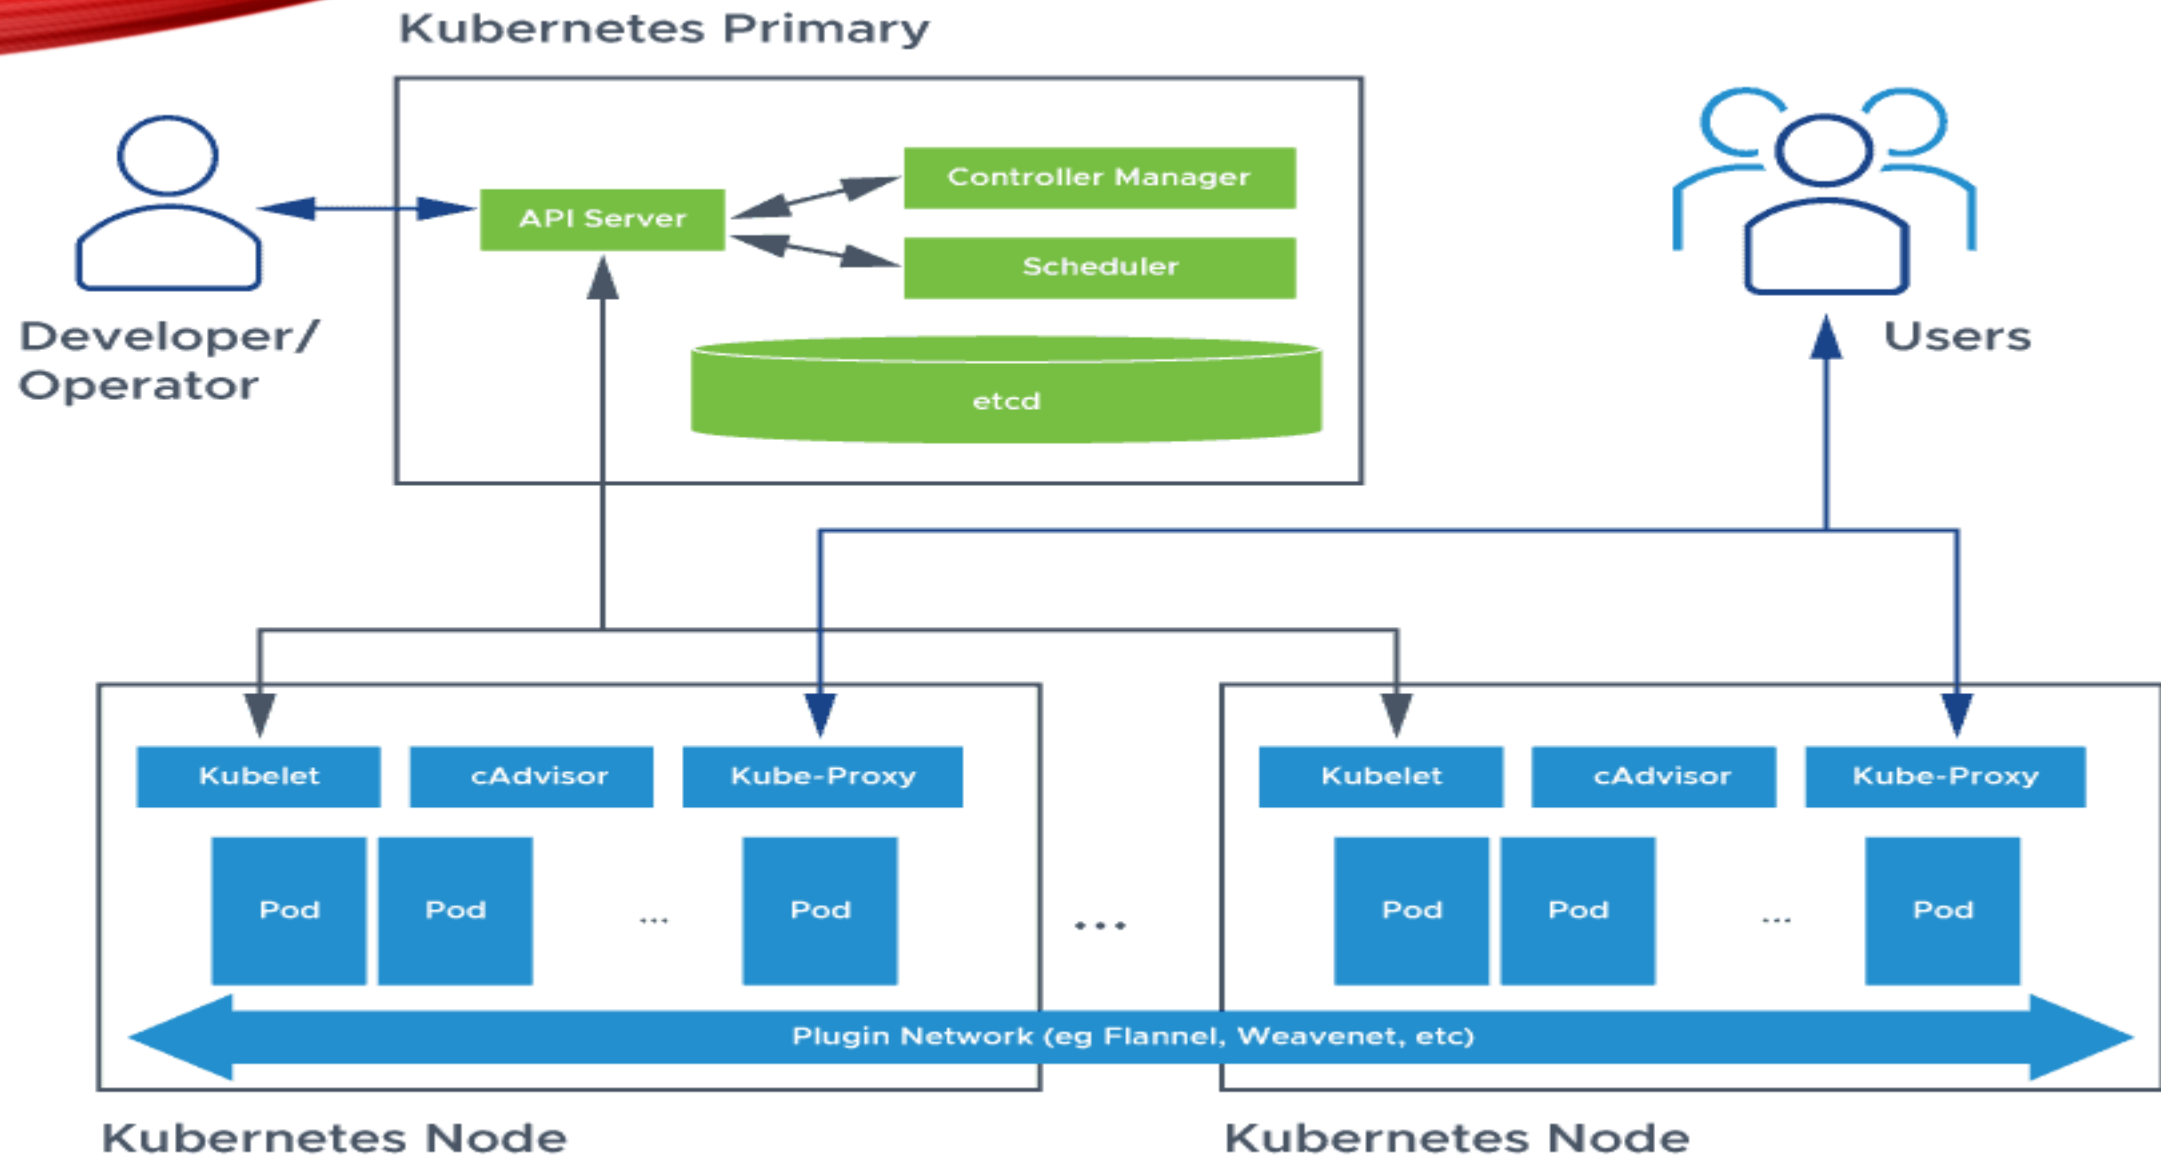
Kubernetes

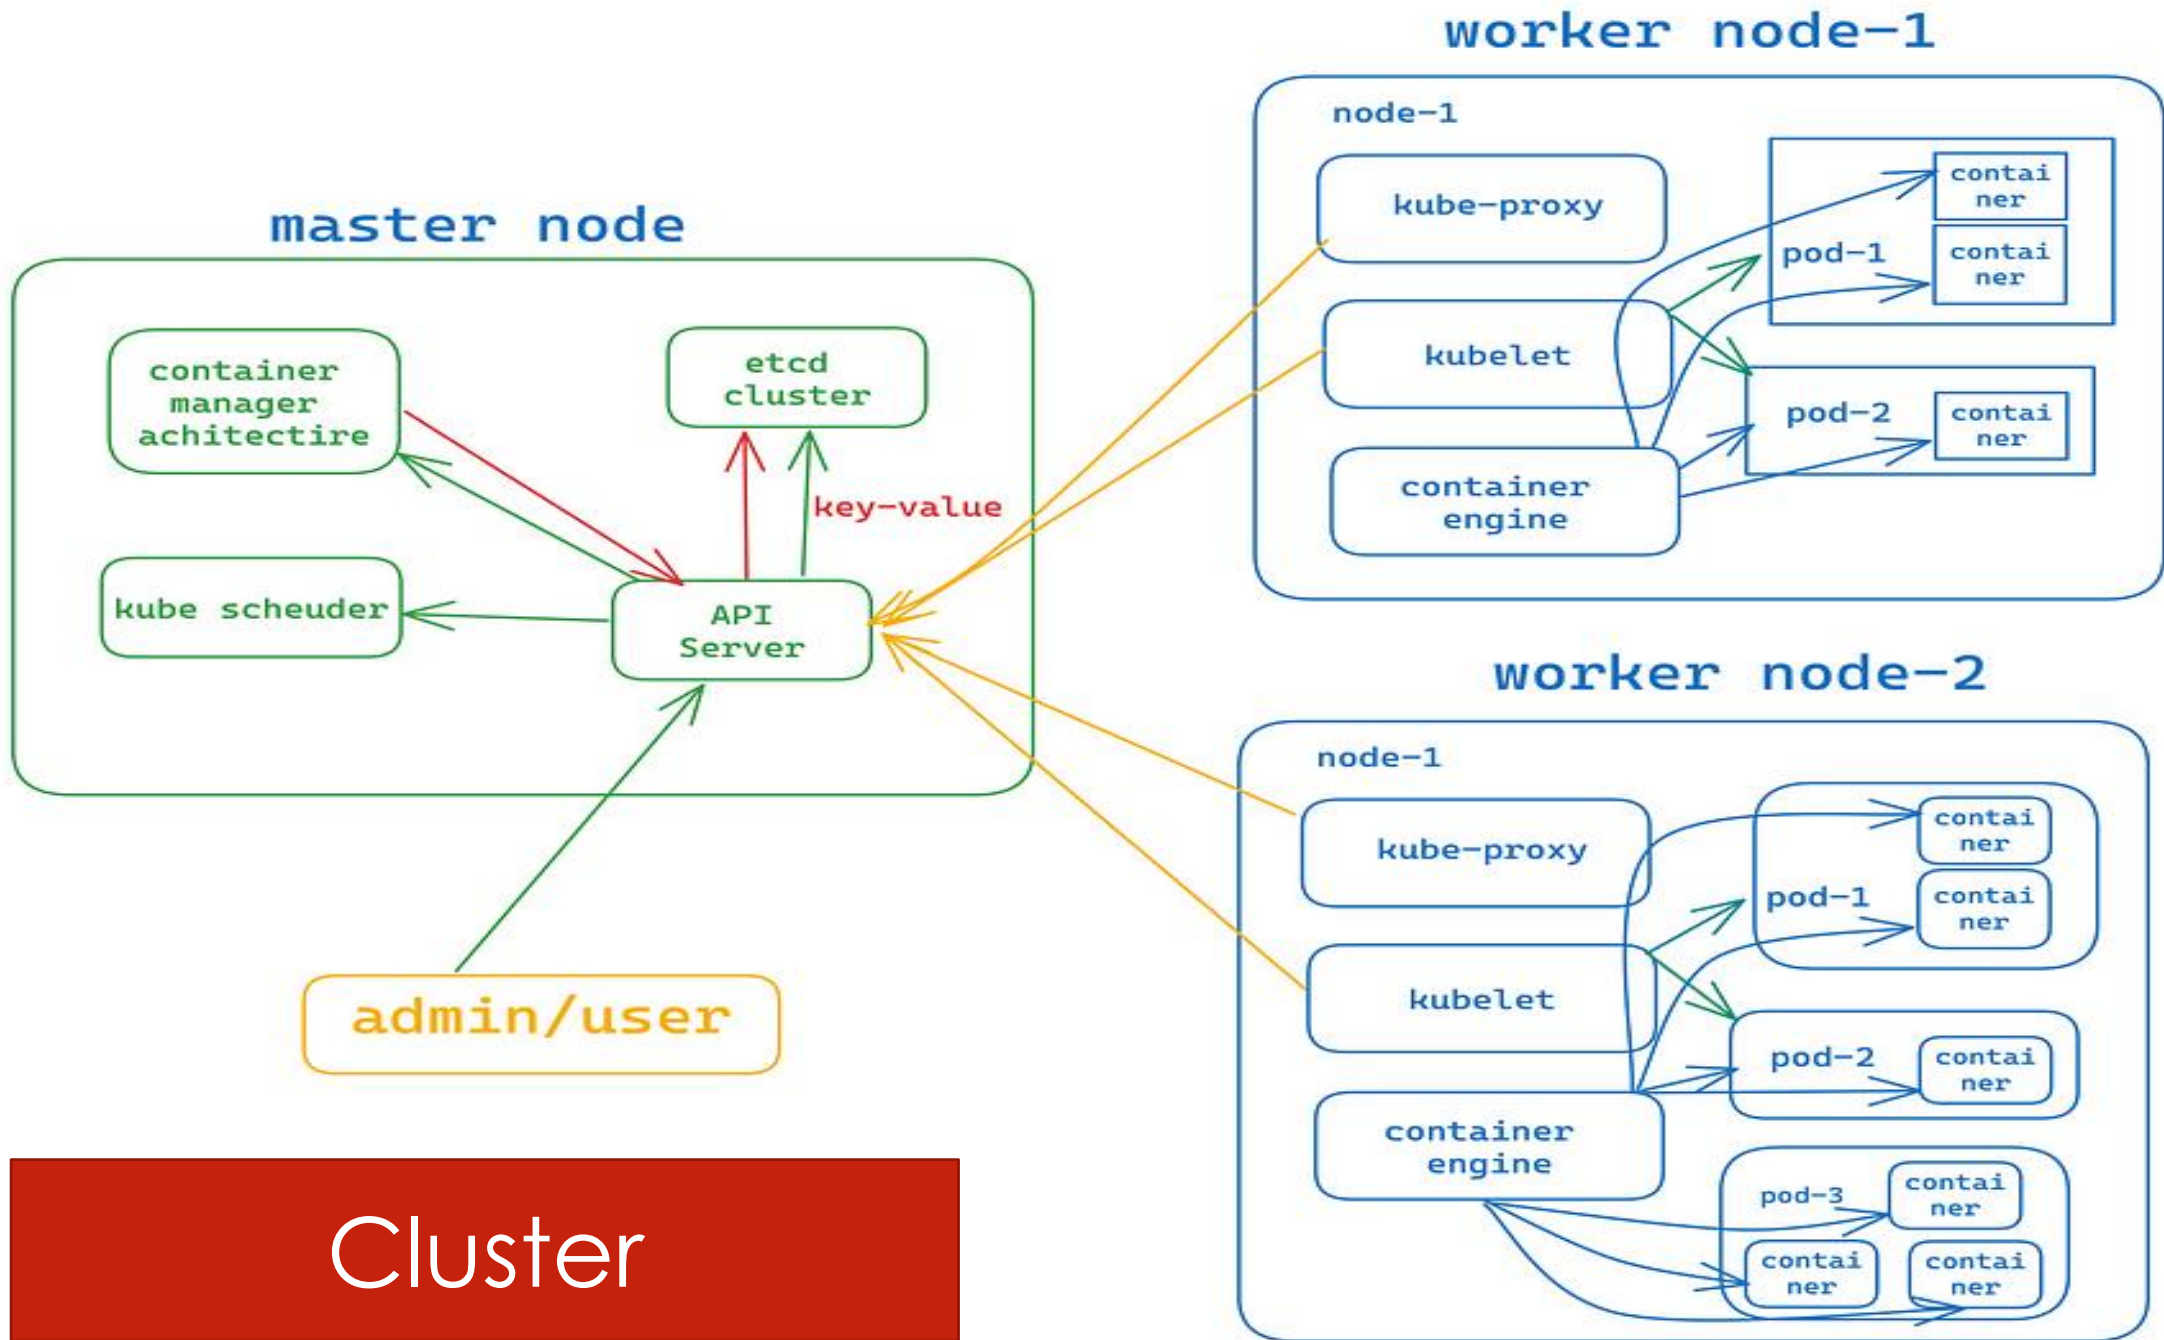
KUBERNETES



- Kubernetes, an open-source container orchestration platform, is designed to automate the deployment, scaling, and operation of application containers.

KUBERNETES ARCHITECTURE






Cluster

KEY COMPONENTS

Master Node: Controls the cluster, responsible for the management of the Kubernetes cluster.

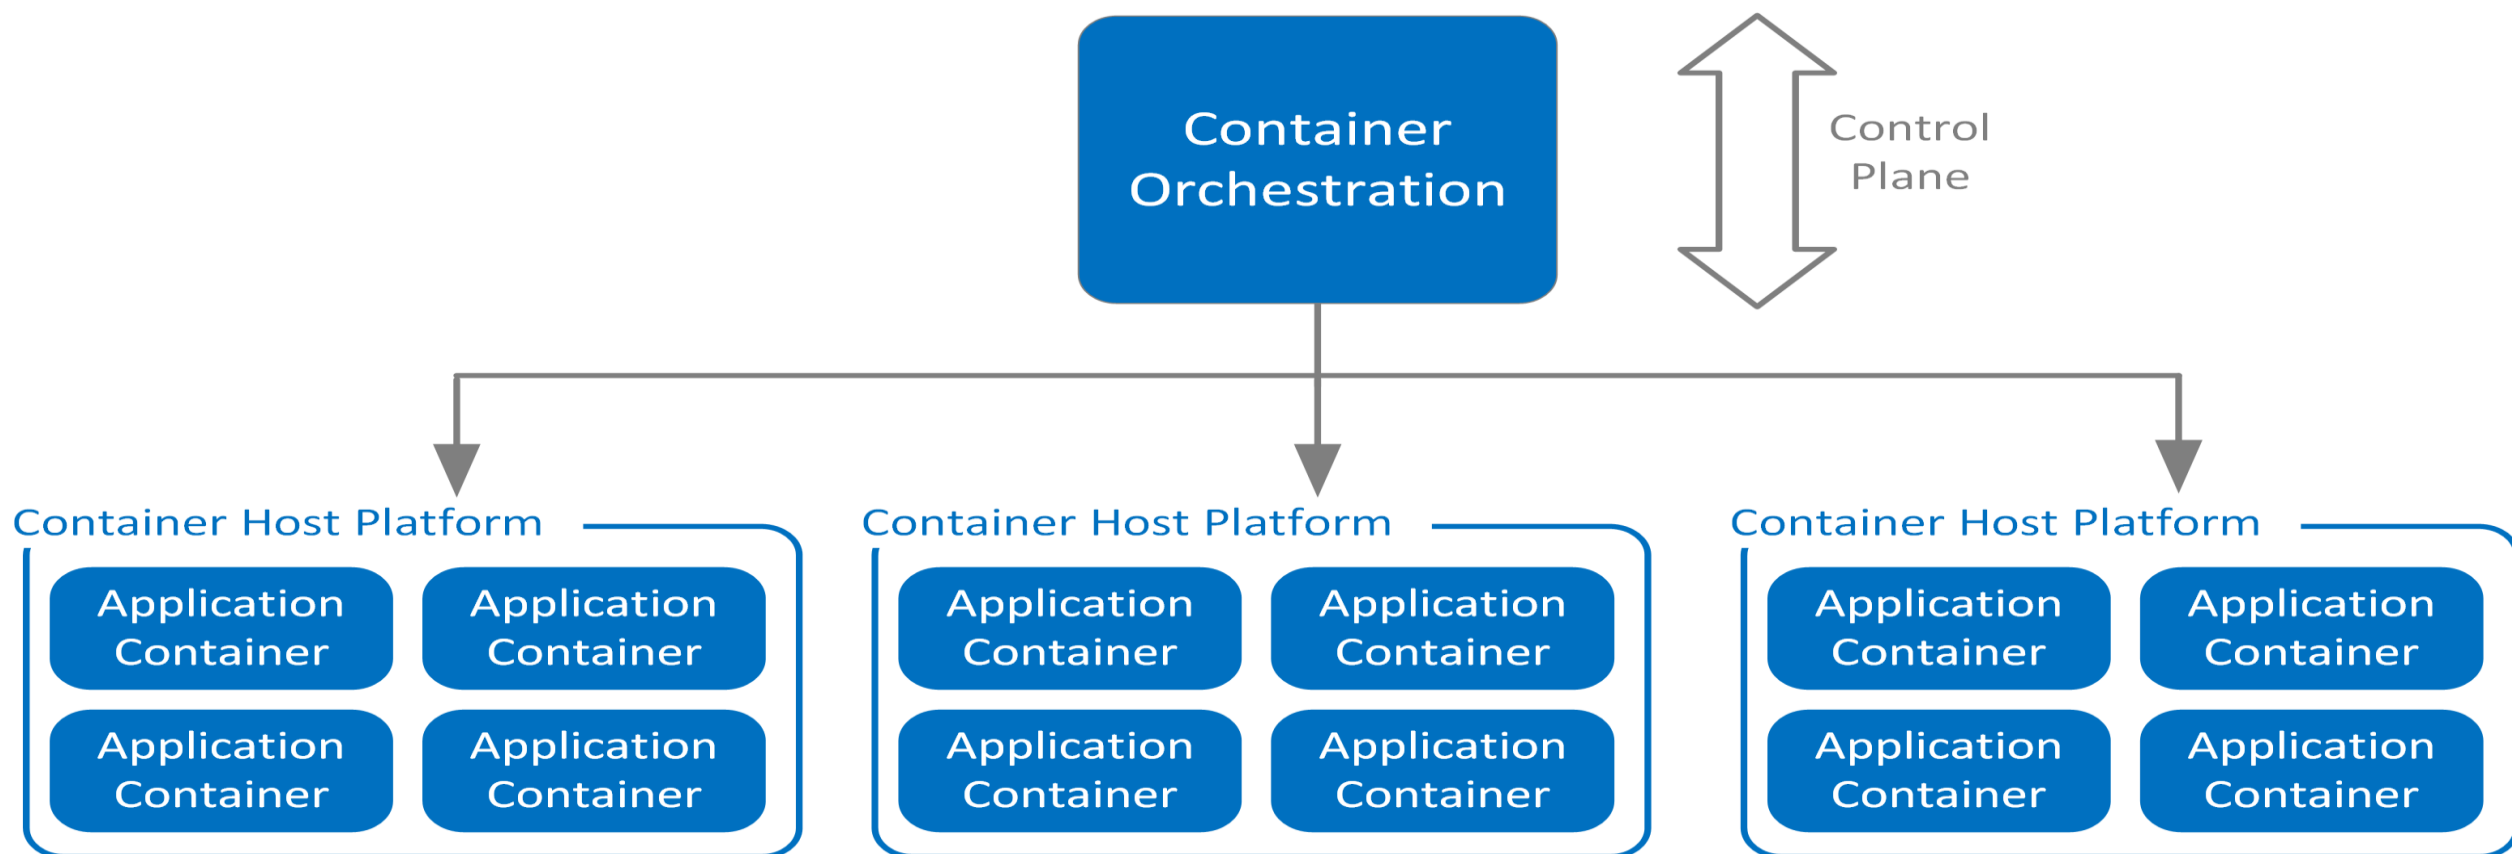
- **API Server:** The front-end for the Kubernetes control plane. It exposes the Kubernetes API, acting as the main management point for the cluster.
- **Scheduler:** Assigns workloads to specific nodes based on resource availability and policies.
- **Controller Manager:** Runs controller processes to regulate the state of the cluster, ensuring that the desired state matches the current state.
- **etcd:** A consistent and highly-available key-value store used as Kubernetes' backing store for all cluster data.



Worker Nodes:
Run containerized applications.
Each node has the necessary services to run pods and is managed by the master node.

- **Kubelet:** An agent that ensures containers are running in a pod. It communicates with the master node and manages the containers' lifecycle.
- **Kube-proxy:** Maintains network rules on nodes, enabling communication to and from pods within the cluster.
- **Container Runtime:** manages and runs components required to run containers
- **Pod:** Represents a single instance of an application

CONTAINER ORCHESTRATION





KUBERNETES API

- The Kubernetes API is a powerful interface that allows users to interact with and manage their Kubernetes clusters.
- It provides a RESTful interface for querying and manipulating the state of various Kubernetes objects, such as pods, services, deployments, and more.

KEY FEATURES OF KUBERNETES API

- **Resource-Oriented Design:**

- The Kubernetes API is structured around resources, such as Pods, Services, Nodes, ConfigMaps, and more.
- Each resource has a corresponding URL endpoint, allowing for operations like create, read, update, and delete (CRUD).

- **RESTful Endpoints:**

- The API uses standard HTTP verbs: GET (retrieve), POST (create), PUT (update), PATCH (partially update), DELETE (remove).
- URLs are structured in a hierarchical format, e.g., `/api/v1/namespaces/{namespace}/pods`.



KEY FEATURES OF KUBERNETES API

- **API Versions:**

- The Kubernetes API is versioned to maintain compatibility. Common versions include v1, v1beta1, and v1alpha1.
- Different API groups (e.g., core, apps, batch) may have different versions.

COMMON API RESOURCES

- **Pods:**

- Smallest and simplest Kubernetes object, representing a single instance of a running process.

- **Services:**

- Abstract way to expose an application running on a set of Pods as a network service.

- **Deployments:**

- Provide declarative updates to applications, managing ReplicaSets to ensure the desired number of Pods are running.

COMMON API RESOURCES

- **ConfigMaps:**

- Store configuration data in key-value pairs.

- **Secrets:**

- Store sensitive data, such as passwords, OAuth tokens, and SSH keys.

- **Nodes:**

- Represent a worker node in the cluster.



INTERACTING WITH THE API

- **kubectl:**
 - The primary command-line tool for interacting with the Kubernetes API.
 - Commands like `kubectl get`, `kubectl create`, `kubectl delete`, etc., interact with the API to manage resources.



INTERACTING WITH THE API

- **Client Libraries:**

- Kubernetes provides client libraries for different programming languages (e.g., Go, Python, Java, JavaScript).
- These libraries wrap the RESTful API calls, making it easier to programmatically interact with Kubernetes.

- **Direct HTTP Calls:**

- Users can make direct HTTP requests to the API server endpoints.
- Authentication is required, typically using tokens, certificates, or other supported methods.

CREATING AND MANAGING CLUSTER

- Create clustuer using Minikube
- Download minikube installer
- Install
- Once instalation done.


minikube version

minikube start

Thi command will start kubernetes cluster.

CHECK BASIC COMMANDS


- On the master node, enter the following command to review cluster information:
 - `kubectl cluster-info`
- On the master node, enter the following command to view complete cluster information
 - `kubectl cluster-info dump`
- Use the following command to create a namespace:
 - `kubectl create namespace firstnamespace`

- 
- Confirm the creation of the new namespace with the following command:
 - `kubectl get namespaces`
 - To view the cluster configuration, use the command below:
 - `kubectl config view`
 - Run the following command to view the current cluster:
 - `kubectl config current-context`
 - To identify the API server, execute and copy the 127.0.0.1:8080 port as shown below:
 - `kubectl proxy --port=8080`

CONFIGURING PODS IN CLUSTER

- Create YAML file: vi pod.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
  - name: my-container
    image: nginx:latest
    ports:
    - containerPort: 80
```

- 
- Apply the updated Pod Definition YAML file:
 - `kubectl apply -f my-pod.yaml`
 - Verify the Pod is running:
 - `kubectl get pods`

KUBERNATES SERVICE

- To access the Nginx container running in your AKS cluster, you can expose it using a Kubernetes service.
- Create a file named my-service.yaml

```
PS /home/sfj> cat my-service.yaml
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: LoadBalancer
  selector:
    app: my-nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```


MODIFY THE POD TO MATCH SERVICE

- Update my-pod.yaml to include a label that matches the selector in the service.

```
PS /home/sfj> cat my-pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
  labels:
    app: my-nginx
spec:
  containers:
  - name: my-container
    image: nginx:latest
    ports:
    - containerPort: 80
```

- Apply the updated Pod Definition YAML file
 - `kubectl apply -f my-pod.yaml`
- Deploy the service to expose the pod:
 - `kubectl apply -f my-service.yaml`
- Check service status:
 - `kubectl get services`
- Copy external Ip address and check in browser for running service

```
PS /home/sfj> kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.0.0.1	<none>	443/TCP	38h
my-service	LoadBalancer	10.0.49.202	172.179.99.34	80:31060/TCP	35s
nginx	LoadBalancer	10.0.188.58	172.179.147.5	80:31882/TCP	38h

RUNNING CONTAINERS IN PODS

- A Pod is the smallest and simplest Kubernetes object.
- It represents a single instance of a running process in your cluster.
- Pods are designed to host one or more containers that share the same network namespace, storage, and specifications.
- You run your application containers inside Pods.
- Each Pod has its own IP address, and containers within a Pod can communicate with each other using localhost.
- Pods can be created, destroyed, and recreated by Kubernetes, especially if part of a ReplicaSet, Deployment, or other higher-level abstraction.

RUNNING SERVICES IN PODS

- A Service in Kubernetes is an abstraction that defines a logical set of Pods and a policy by which to access them.
- Services provide stable network endpoints to Pods, allowing you to expose your applications internally or externally.
- Services are used to expose Pods to other Pods within the cluster, or to external users.
- They can load balance traffic across multiple Pods.
- Kubernetes supports several types of Services, including
 - ClusterIP
 - NodePort
 - LoadBalancer

PRACTICAL USE CASE

Pods:

- You deploy an application by creating Pods.
- For example, deploying an Nginx server involves creating a Pod with an Nginx container.

Services:

- You expose the Nginx application to the internet or within the cluster by creating a Service that targets the Nginx Pod.



LET'S SUMMARIZE

- Pods run the actual application containers and manage their lifecycle
- while Services provide stable network endpoints to access those Pods, handle load balancing, and expose the application to other components or external clients.



DEPLOYMENT

- A Deployment in Kubernetes is a higher-level abstraction that manages the lifecycle of Pods and ReplicaSets.


DEPLOYMENT FEATURES

- **Declarative Updates:** You declare the desired state of your application, and the Deployment controller makes the necessary changes to achieve that state.
- **Scaling:** Deployments allow you to scale the number of replicas of your application up or down.
- **Rolling Updates:** You can update your application with zero downtime by rolling out updates incrementally.
- **Rollbacks:** If an update causes issues, you can roll back to a previous stable state.
- **Self-Healing:** If a Pod fails, the Deployment controller replaces it automatically.

CREATING AND CONFIGURING A DEPLOYMENT

- Create a Deployment YAML File:
- Nano deployment.yaml


```
PS /home/sfj> cat deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.17.4
        ports:
        - containerPort: 80
```

- 
- Apply the Deployment:
 - `kubectl apply -f deployment.yaml`
 - Verify the Deployment:
 - `kubectl get deployments`
 - `kubectl get pods`
 - Deployment entire Description
 - `kubectl describe deployment nginx-deployment`
 - Scaling the Deployment:
 - `kubectl scale deployment/nginx-deployment --replicas=5`
 - You can see the 5 replica running in deployment

UPDATE DEPLOYMENT

```
PS /home/sfj> cat deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.18.0 #Updated Image
        ports:
        - containerPort: 80
```

- changing the container image
- modify the YAML file: nano deployment.yaml

- 
- Apply the updated YAML file:
 - `kubectl apply -f deployment.yaml`
 - Rolling Back a Deployment:
 - `kubectl rollout undo deployment/nginx-deployment`
 - Monitoring the deployment
 - `kubectl rollout status deployment/nginx-deployment`
 - Check the deployment description which will show the version rollbacked
 - `kubectl describe deployment nginx-deployment`